

# CPSC429/529: Machine Learning

## Lecture 1: Introduction

Dongsheng Che  
Computer Science Department  
East Stroudsburg University



### **Acknowledgment:**

Some slides adapted from Kevyn Collins-Thompson's Applied ML class on [Coursera](#).

# **Course Syllabus**

# Course Information

- Textbooks:
  - Required: Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems by Aurélien Géron (AG book).
  - Optional: Data Science from Scratch: First Principles with Python, by Joel Grus



# Course Prerequisites

- CPSC 380: Introduction to Data Science

- Python language



- Numpy: Scientific Computing Library



- Pandas: Data Manipulation and Analysis

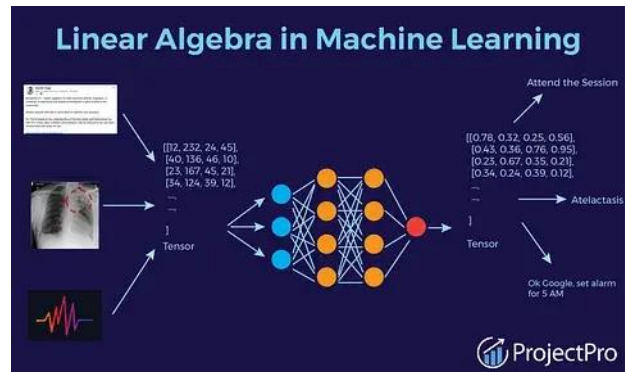


- Matplotlib: Visualization library



# Other skills you need

- Simple linear algebra (vectors, matrices): Math 320



- Basic probability theory: Math 311

$$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$$

Probability of B occurring given evidence A has already occurred

Probability of A occurring

Probability of A occurring given evidence B has already occurred

Probability of B occurring

# What if I don't have these prerequisites?

Options:

- Take CPSC 380 (Introduction to Data Science)
- Stay in the class (If you are a faster learner and can teach numpy, pandas, matplotlib by yourself within a short period of time)

# Class Topics

- Introduction to Machine Learning
- Scikit-learn: Machine Learning Framework
  - Predictive model pipeline
  - Select the best model
  - Hyperparameter tuning
  - Model evaluations
- Supervised learning:
  - Tree models and model ensembles
  - Linear models: Linear regression, logistic regression
- Neural networks and deep learning
  - Introduction to ANN with Keras
  - Training Deep NN
  - Custom Models and Training with TensorFlow
- Unsupervised learning:
  - Clustering, Dimensionality reduction

# Grading Policy

## **CPSC429:**

- Labs (25 pts)
- Programming assignments (20 pts)
- Term project (20 pts)
- Midterm exam (35pts)

## **CPSC529:**

- Labs (20 pts)
- Programming assignments (20 pts)
- Term project (20 pts)
- Midterm exam (30pts)
- Literature Research Paper (10 pts)



# **Course Syllabus**

# Introduction to ML

What is Machine  
Learning?

# ARTIFICIAL INTELLIGENCE VS MACHINE LEARNING VS DEEP LEARNING

## 1 Artificial Intelligence

Development of smart systems and machines that can carry out tasks that typically require human intelligence


## 2 Machine Learning

Creates algorithms that can learn from data and make decisions based on patterns observed  
Require human intervention when decision is incorrect

## 3 Deep Learning

Uses an artificial neural network to reach accurate conclusions without human intervention

# Applications of Machine learning



# What is Machine Learning (ML)?

- **The study of computer programs (algorithms) that can learn by example**
- **ML algorithms can generalize from existing examples of a task**
  - *e.g. after seeing a training set of labeled images, an image classifier can figure out how to apply labels accurately to new, previously unseen images*

# Machine Learning models can learn by example

- Algorithms learn rules from labelled examples
- A set of labelled examples used for learning is called training data.
- The learned rules should also be able to generalize to correctly recognize or predict new examples not in the training set.

Audio signal



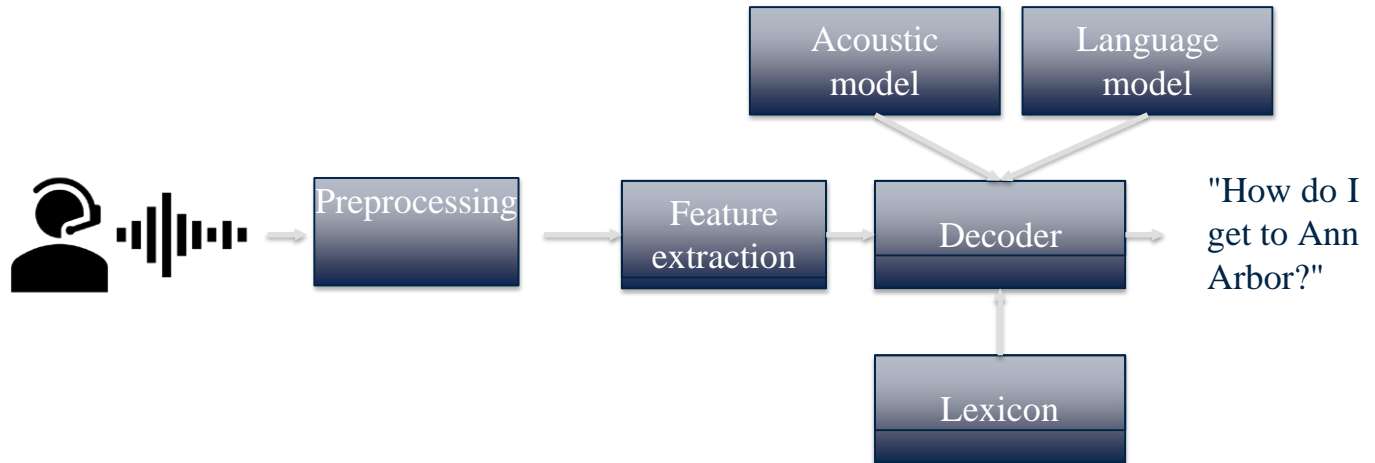
Output text

How do I  
get to Ann  
Arbor?

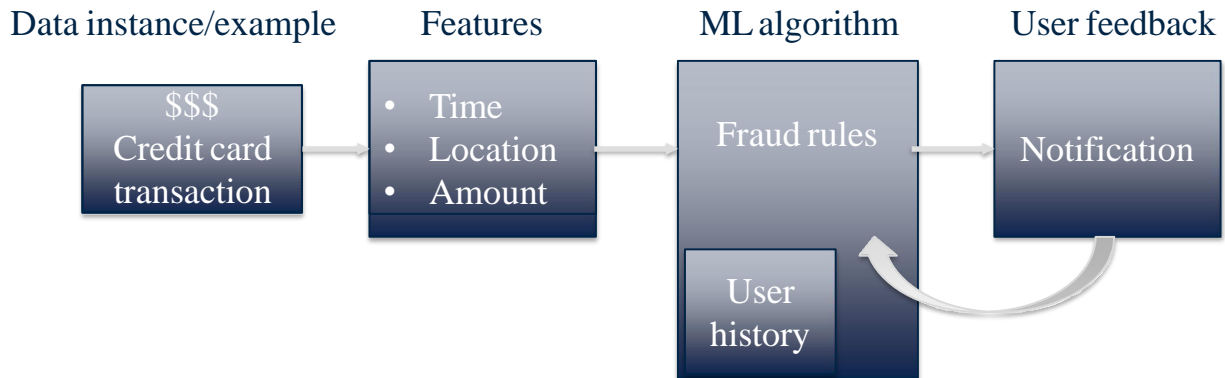
Hello!

Please order  
me a pizza.

# Machine Learning for Speech Recognition



# Machine Learning for fraud detection and credit scoring

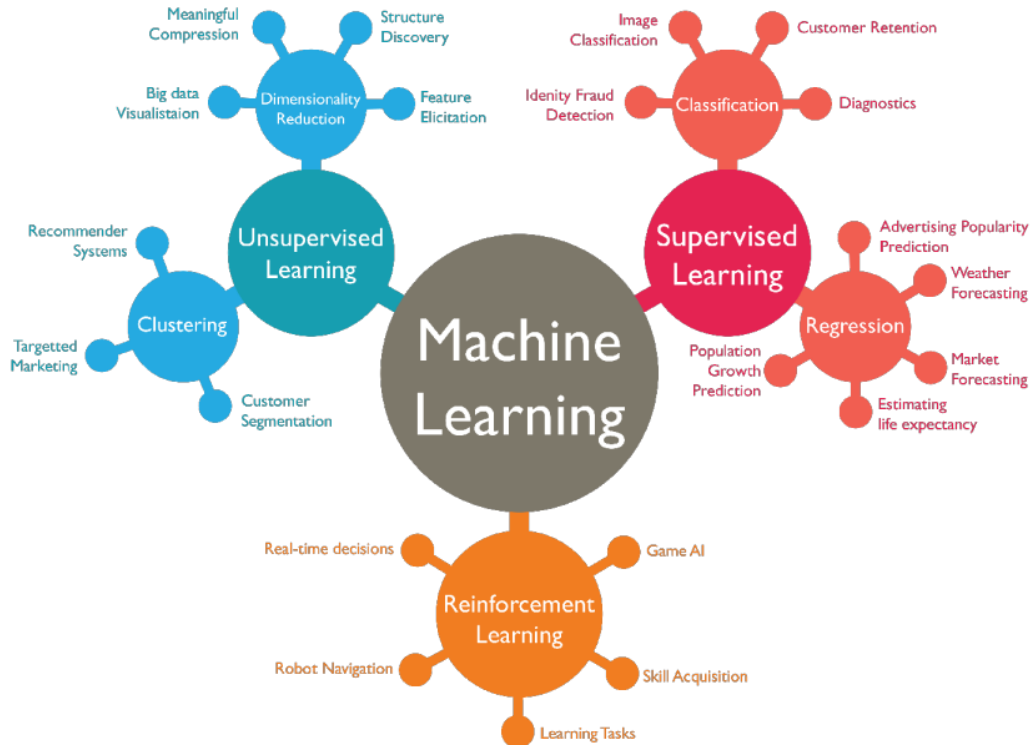




# **Introduction to ML**

Types of Machine Learning  
problems

# Types of Machine Learning Models







# Key types of Machine Learning problems

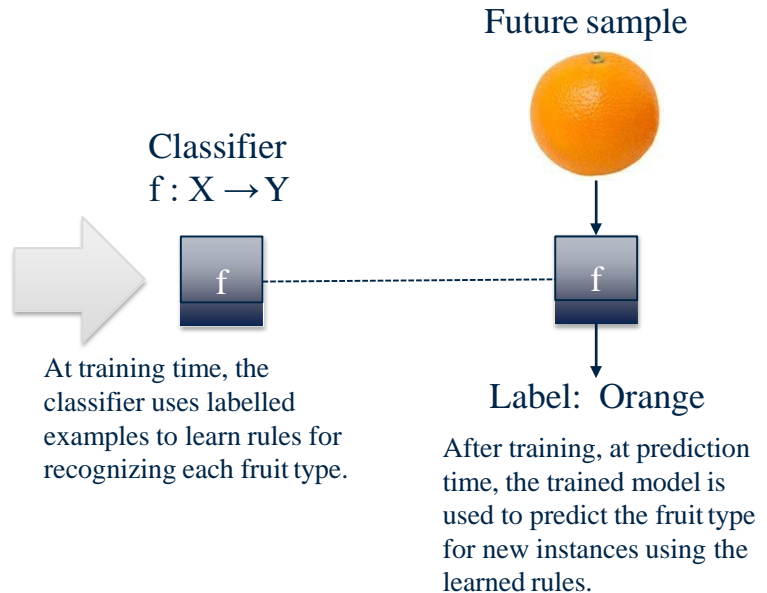
Supervised machine learning: Learn to predict target values from labelled data.

- **Classification** (target values are discrete classes)
- **Regression** (target values are continuous values)

# Supervised Learning (classification example)

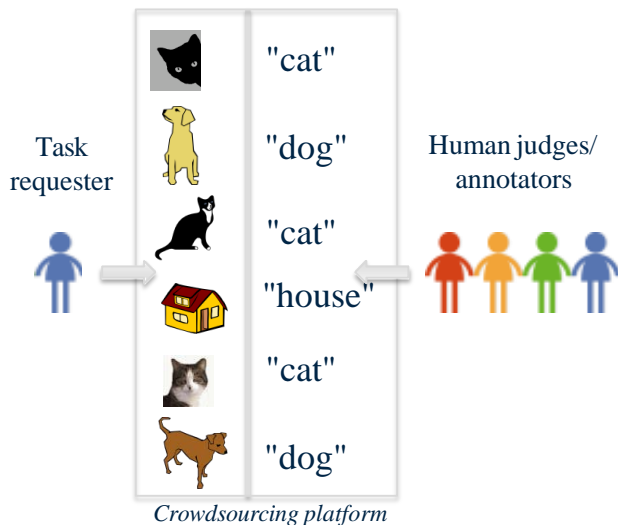
Training set

X Sample	Y Target Value (Label)
 $xx_1$	Apple $\mathcal{Y}_1$
 $xx_2$	Lemon $\mathcal{Y}_2$
 $xx_3$	Apple $\mathcal{Y}_3$
 $xx_4$	Orange $\mathcal{Y}_4$

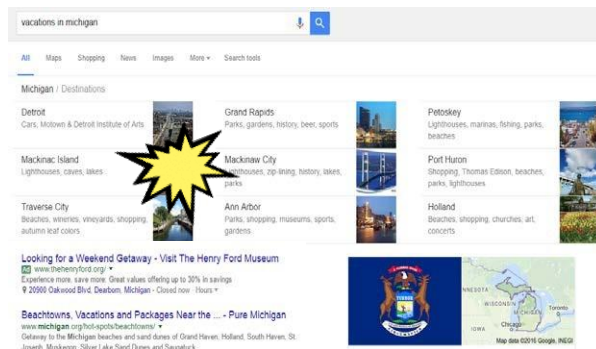


# Examples of explicit and implicit label sources

## Explicit labels



## Implicit labels



Clicking and reading the "Mackinac Island" result can be an implicit label for the search engine to learn that "Mackinac Island" is especially relevant for the query [vacations in michigan] for that specific user.

# Key types of Machine Learning problems

Supervised machine learning: Learn to predict target values from labelled data.



- Classification (target values are discrete classes)
- Regression (target values are continuous values)

Unsupervised machine learning: Find structure in *unlabeled data*

- Find groups of similar instances in the data (clustering)
- Finding unusual patterns (outlier detection)

# Unsupervised learning: finding useful structure or knowledge in data when no labels are available

- Finding clusters of similar users (clustering)
- Detecting abnormal server access patterns (unsupervised outlier detection)



# **Introduction to ML**

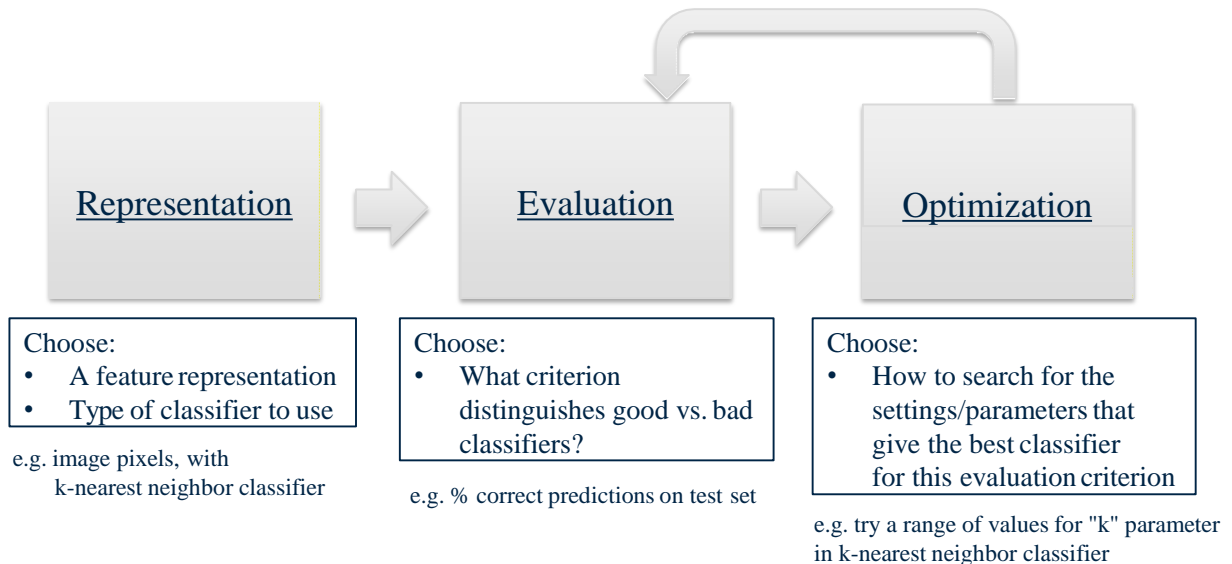
Types of Machine Learning  
problems



# **Introduction to ML**

Machine Learning Workflow

# A Basic Machine Learning Workflow



# Feature Representations

## Email

To: Chris Brooks  
From: Daniel Romero  
Subject: Next course offering  
Hi Daniel,  
Could you please send the outline for the  
next course offering? Thanks! -- Chris



<u>Feature</u>	<u>Count</u>
to	1
chris	2
brooks	1
from	1
daniel	2
romero	1
the	2
...	

## Feature representation

A list of words with  
their frequency counts

## Picture



A matrix of color  
values (pixels)


## Sea Creatures



<u>Feature</u>	<u>Value</u>
DorsalFin	Yes
MainColor	Orange
Stripes	Yes
StripeColor1	White
StripeColor2	Black
Length	4.3 cm

A set of attribute values

# Representing a piece of fruit as an array of features (plus label information)



## 1. Feature representation

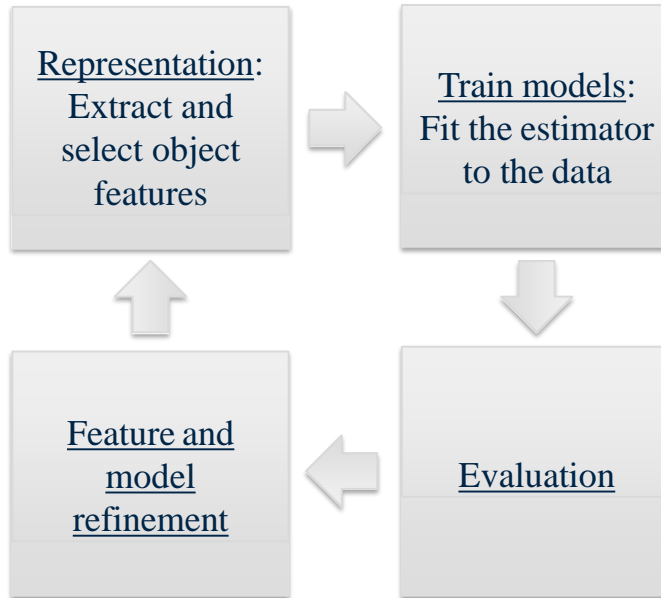
Label information (available in training data only)				Feature representation			
	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
18	1	apple	cripps_pink	162	7.5	7.1	0.83

## 2. Learning model

Classifier

Predicted class  
(apple)

# Represent / Train / Evaluate / Refine Cycle



# **Introduction to ML**

Machine Learning Workflow

# **Introduction to ML**

Python Tools for Machine  
Learning

# NumPy: Scientific Computing Library




<http://www.numpy.org/>

- Provides fundamental data structures used by scikit-learn, particularly multi-dimensional arrays.
- Typically, data that is input to scikit-learn will be in the form of a NumPy array.
- Example import: `import numpy as np`



# Pandas: Data Manipulation and Analysis

pandas  
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



<http://pandas.pydata.org/>

- Provides key data structures like `DataFrame`
- Also, support for reading/writing data in different formats
- Example import: `import pandas as pd`

# matplotlib and other plotting libraries

**matplotlib**  <http://matplotlib.org/>

- We typically use matplotlib's **pyplot** module:  

```
import matplotlib.pyplot as plt
```
- We also sometimes use the **seaborn** visualization library (<http://seaborn.pydata.org/>)  

```
import seaborn as sn
```
- And sometimes the **graphviz** plotting library:  

```
import graphviz
```

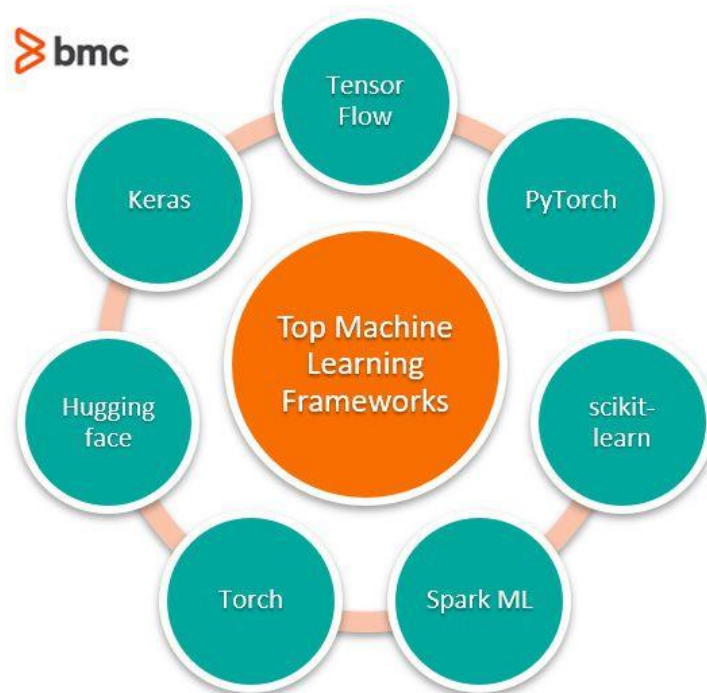
# SciPy Library: Scientific Computing Tools



<http://www.scipy.org/>

- Provides a variety of useful scientific computing tools, including statistical distributions, optimization of functions, linear algebra, and a variety of specialized mathematical functions.
- With scikit-learn, it provides support for *sparse matrices*, a way to store large tables that consist mostly of zeros.
- Example import: `import scipy as sp`

# Top Machine Learning Frameworks



# scikit-learn: Python Machine Learning Library



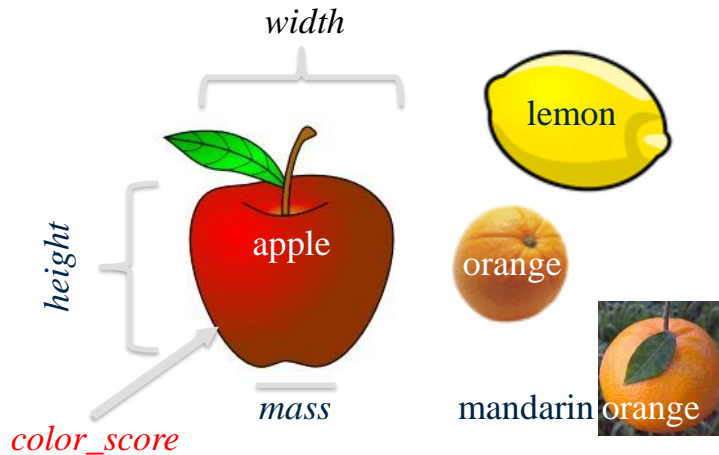
- **scikit-learn Homepage**  
<http://scikit-learn.org/>
- **scikit-learn User Guide**  
[http://scikit-learn.org/stable/user\\_guide.html](http://scikit-learn.org/stable/user_guide.html)
- **scikit-learn API reference**  
<http://scikit-learn.org/stable/modules/classes.html>
- **In Python, we typically import classes and functions we need like this:**  

```
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier
```

# **Introduction to ML**

An Example Machine  
Learning Problem

# The Fruit Dataset



	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67

fruit\_data\_with\_colors.txt

Credit: Original version of the fruit dataset created by Dr. Iain Murray, Univ. of Edinburgh

# The input data as a table

Each row corresponds to a single data instance (sample)


	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	190	8.0	6.9	0.56
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67
17	1	apple	golden_delicious	168	7.5	7.6	0.73
18	1	apple	cripps_pink	162	7.5	7.1	0.83
19	1	apple	cripps_pink	162	7.4	7.2	0.85
20	1	apple	cripps_pink	160	7.5	7.5	0.86

The fruit\_label column contains the label for each data instance (sample)

These four columns contain the features of each data instance (sample)

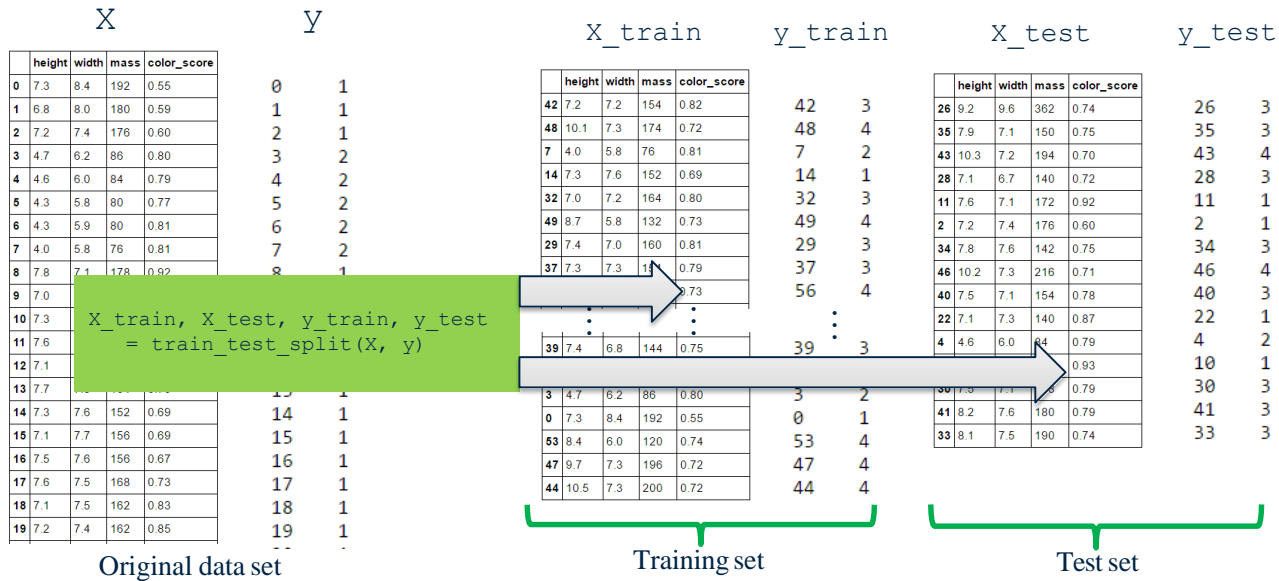


# The scale for the (simplistic) `color_score` feature used in the fruit dataset



Color category	<code>color_score</code>
Red	0.85 - 1.00
Orange	0.75 - 0.85
Yellow	0.65 - 0.75
Green	0.45 - 0.65

# Creating Training and Testing Sets




# Some reasons why looking at the data initially is important

- Inspecting feature values may help identify what **cleaning or preprocessing** still needs to be done once you can see the range or distribution of values that is typical for each attribute.
- You might notice **missing or noisy data, or inconsistencies** such as the wrong data type being used for a column, incorrect units of measurements for a particular column, or that there aren't enough examples of a particular class.
- You may realize that your problem is actually solvable without machine learning.

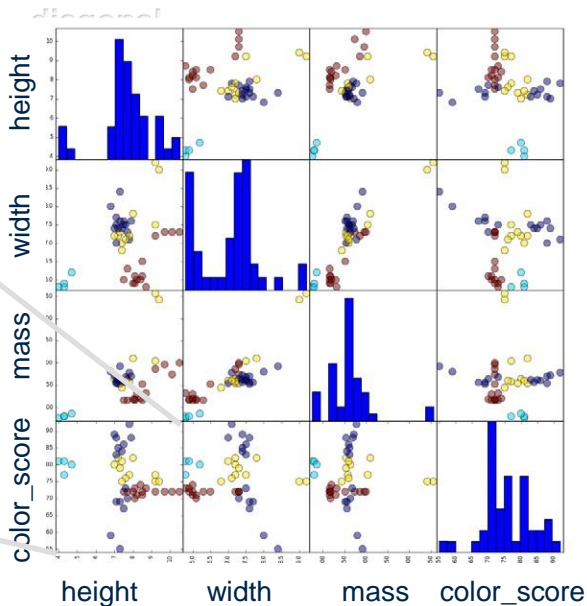
Examples of incorrect or missing feature values

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	192
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	apple	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn		7.4	7.0	0.89
10	1	apple	braeburn		6.9	7.3	0.93
11	1	apple	braeburn		7.1	7.6	0.92
12	1	apple	braeburn		7.0	7.1	0.88
13	1	apple	golden_delicious		7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69

**A pairwise feature scatterplot visualizes the data using all possible pairs of features, with one scatterplot per feature pair, and histograms for each feature along the diagonal**



Individual scatterplot plotting all fruits by their **height** and **color\_score**.  
Colors represent different fruit classes.



```
In [27]: fruits
```

```
Out[27]:
```

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67

```
In [88]: fruits.shape
```

```
Out[88]: (59, 7)
```

In [88]: fruits.shape

Out[88]: (59, 7)

In [92]: X\_train.shape

Out[92]: (44, 4)

In [93]: X\_test.shape

Out[93]: (15, 4)

In [94]: y\_train.shape

Out[94]: (44,)

In [95]: y\_test.shape

Out[95]: (15,)

In [96]: X\_train

Out[96]:

	height	width	mass	color_score
42	7.2	7.2	154	0.82
48	10.1	7.3	174	0.72
7	4.0	5.8	76	0.81
14	7.3	7.6	152	0.69
32	7.0	7.2	164	0.80
49	8.7	5.8	132	0.73
29	7.4	7.0	160	0.81
37	7.3	7.3	154	0.79
56	8.1	5.9	116	0.73
18	7.1	7.5	162	0.83
55	7.7	6.3	116	0.72
27	9.2	7.5	204	0.77
15	7.1	7.7	156	0.69
5	4.3	5.8	80	0.77
31	8.0	7.8	210	0.82
16	7.5	7.6	156	0.67

In [98]: y\_train

Out[98]:

42	3
48	4
7	2
14	1
32	3
49	4
29	3
37	3
56	4
18	1
55	4
27	3
15	1
5	2
31	3
16	1
50	4
20	1
51	4
8	1
13	1
25	3
17	1
58	4
57	4
52	4
38	3
1	1
12	1
45	4
24	3
6	2

In [97]: X\_test

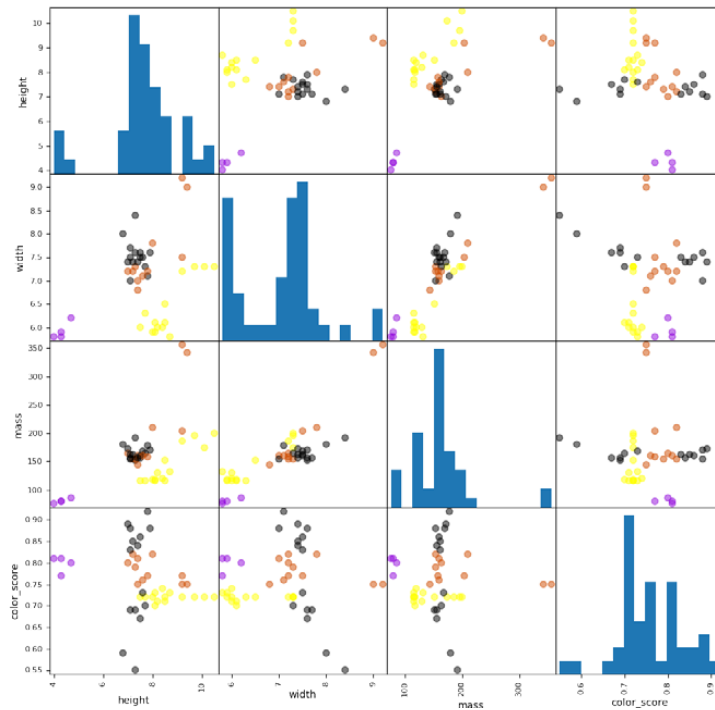
Out[97]:

	height	width	mass	color_score
26	9.2	9.6	362	0.74
35	7.9	7.1	150	0.75
43	10.3	7.2	194	0.70
28	7.1	6.7	140	0.72
11	7.6	7.1	172	0.92
2	7.2	7.4	176	0.60
34	7.8	7.6	142	0.75
46	10.2	7.3	216	0.71
40	7.5	7.1	154	0.78
22	7.1	7.3	140	0.87
4	4.6	6.0	84	0.79
10	7.3	6.9	166	0.93
30	7.5	7.1	158	0.79
41	8.2	7.6	180	0.79
33	8.1	7.5	190	0.74

In [99]: y\_test

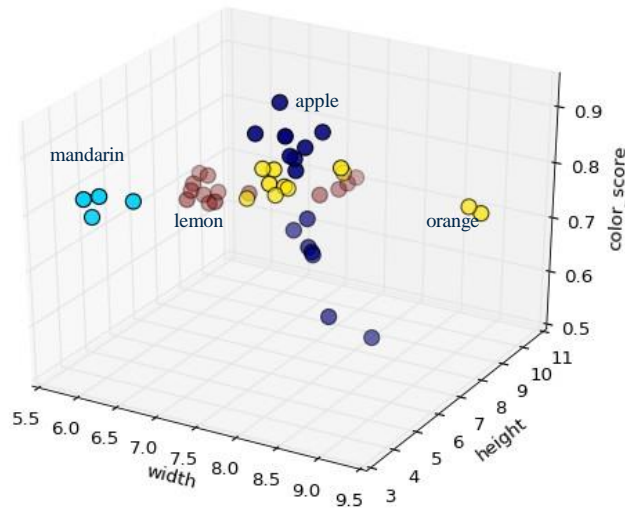
Out[99]:

26	3
35	3
43	4
28	3
11	1
2	1
34	3
46	4
40	3
22	1
4	2
10	1
30	3
41	3
33	3



```
from matplotlib import cm
cmap = cm.get_cmap('gnuplot')
scatter = pd.scatter_matrix(X_train, c= y_train, marker = 'o', s=40, hist_kws={'bins':15}, figsize=(12,12), cmap=cmap)
```

# A three-dimensional feature scatterplot



```
from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.add_subplot(111, projection = '3d')
ax.scatter(X_train['width'], X_train['height'], X_train['color_score'], c = y_train, marker = 'o', s=100)
ax.set_xlabel('width')
ax.set_ylabel('height')
ax.set_zlabel('color_score')
plt.show()
```



# **Introduction to ML**

An Example Machine  
Learning Problem


# **Introduction to ML**

K-Nearest Neighbors  
Classification

# The k-Nearest Neighbor (k-NN) Algorithm

Given a training set  $X_{\text{train}}$  with labels  $y_{\text{train}}$ , and given a new instance  $x_{\text{test}}$  to be classified:

1. Find the most similar instances (let's call them  $X_{\text{NN}}$ ) to  $x_{\text{test}}$  that are in  $X_{\text{train}}$ .
2. Get the labels  $y_{\text{NN}}$  for the instances in  $X_{\text{NN}}$
3. Predict the label for  $x_{\text{test}}$  by combining the labels  $y_{\text{NN}}$ . e.g. simple majority vote



# **A nearest neighbor algorithm needs four things specified**

1. A distance metric
2. How many 'nearest' neighbors to look at?
3. Optional weighting function on the neighbor points
4. Method for aggregating the classes of neighbor points

# A nearest neighbor algorithm needs four things specified

1. A distance metric

**Typically Euclidean (Minkowski with  $p = 2$ )**

2. How many 'nearest' neighbors to look at?

**e.g. five**

3. Optional weighting function on the neighbor points


**Ignored**

4. How to aggregate the classes of neighbor points


Simple **majority vote**

(Class with the most representatives among nearest neighbors)

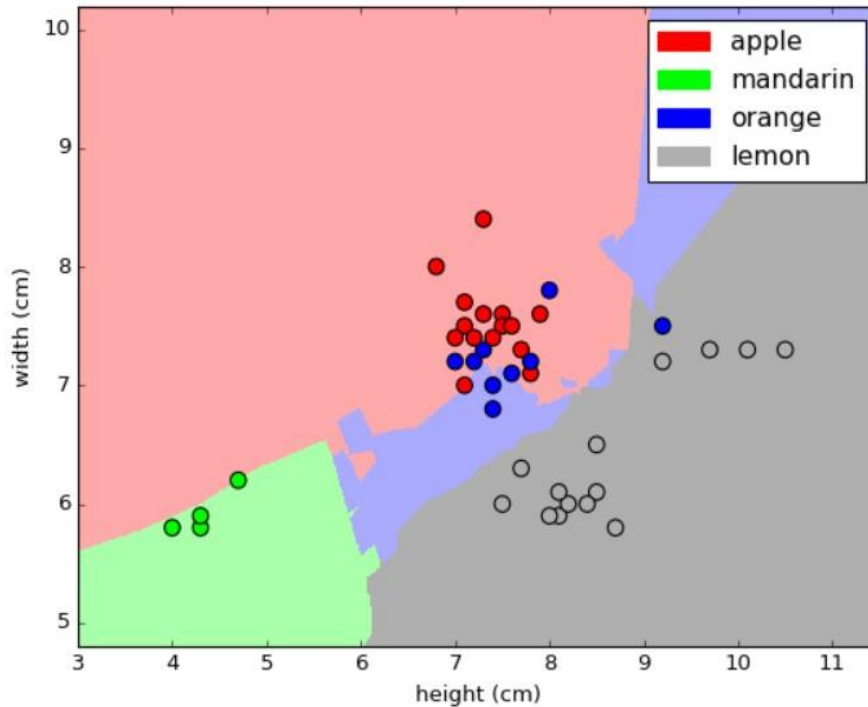
# K-nearest neighbors (k=1) for fruit dataset




# K-nearest neighbors (k=5) for fruit dataset




# K-nearest neighbors (k=10) for fruit dataset








K=1




K=5



K=10

## How sensitive is k-NN classifier accuracy to the choice of 'k' parameter?



Fruit dataset  
with 75%/25%  
train-test split

# Introduction to ML

K-Nearest Neighbors  
Classification